CS 11 Data Structures and Algorithms

Lesson 7: Functions 2

Skip to Main Content

Section 7.1: Pass-By-Value Versus Pass-By-Reference

Many students struggle to understand the difference between the pass-by-value and pass-by-reference parameter passing mechanisms and when to use each. In section 1 of this lesson I will list 7 strategies for describing the difference between the two. My hope is that those of you who seem to be getting the hang of this will gain an even deeper understanding of the issue by hearing it said in many different ways, and those of you who are struggling will find at least one strategy for describing the difference that makes sense to you.

A quick preliminary note: The choice between pass-by-value and pass-by-reference is made for each individual parameter, not just once for each function. So it is possible to have a function with one or more pass-by-value parameters in addition to one or more pass-by-reference parameters.

Strategy #1) The first strategy for describing the difference is the most practical. You can actually apply this without really understanding pass-by-value and pass-by-reference. The first half of the rule is very easy: **if the value of the parameter doesn't change inside the function, use pass-by-value**. There are three ways to change a value inside a function: an assignment statement with the parameter on the left side of the assignment operator, an extraction operator with the parameter on the right side, or a function call with the parameter as an argument. If none of these three situations appears in the function, then you know that you should use pass-by-value. The second half of this first rule is a little more complex: **if the value of the parameter changes, AND you want the calling function to know about the change, use pass-by-reference**.

Strategy #2) Determine the direction that information is flowing. If the information is flowing from the calling function to the called function, use pass-by-value. If the information is flowing from the called function to the calling function, use pass-by-reference. If the information is flowing in both directions, use pass-by-reference. This is essentially the strategy that I used in explaining the difference originally (in the last lesson).

Strategy #3) Look at the situation from the perspective of the called function. If the information is flowing **in** to the function, use pass-by-value. If the information is flowing **out** of the function, use pass-by-reference. If the information is flowing in both directions, use pass-by-reference. There are a couple of pitfalls with this strategy that I should warn you about. What if the called function sends the parameter down to a third function? Does that make the parameter an "out" parameter? The answer is no. When determining whether a parameter is an "in" parameter or an "out" parameter, you should only consider the interaction between the function in question and the function that called it. This possible confusion is one reason why I choose to describe the difference using Strategy #2 rather than this strategy.

Strategy #4) If the function needs to pass information to the function that called it, use pass-by-reference. Otherwise, use pass-by-value. Okay, now we are starting to say the same thing with just slightly different words. But maybe this will click for someone out there.

Strategy #5) If you use pass-by-reference, the changes made to the parameter inside this function are reflected in the calling function. If you use pass-by-value, they are not.

Strategy #6) In a structure diagram, if the arrows are going down use pass-by-value. If the arrows are going up, use pass-by-reference. A structure diagram is a picture you draw of your program. Each function, including main, is represented with a box. You put main at the top, and each function that main calls is placed below main and connected to main with a line. Then you repeat this process for each of the other functions. Finally, you draw arrows on the diagram for each parameter that is passed. For example, in our

draw box program, there would be an arrow starting at getDimensions and going up to main representing the fact that getDimensions is passing the value "width" up to main.

Strategy #7) When you use pass-by-value, the called function receives a value. When you use pass-by-reference, the called function receives a reference, which is a memory address. This is the real behind the scenes technical explanation of what happens when you pass parameters. More details on this in the next section of this lesson.

Section 7.2: The Real Scoop on Passing Parameters

Note: I apologize for the low tech graphics in this section. I'll fix them up when I have time. Better yet, you could volunteer to create some for me...

In this section we will be discussing what really goes on behind the scenes when you pass a parameter either by value or by reference. Some might argue that this level of understanding is not necessary when you are just learning a programming language for the first time. I have found, however, that most of my students find that their ability to use parameters correctly is enhanced by understanding the mechanisms involved at a deeper level.

Let's start by tracing through the execution of our draw box program in detail. I'll be drawing lots of repetitive pictures during this process, so the lesson may seem longer than it actually is. You'll probably want to be looking at a copy of the final version of our draw box program while you read through this. At each stage of our trace, I will show you a snapshot of the current situation along with the relevant code.

First we execute the two declaration statements inside main. A declaration statement like `int width;` tells C++ to allocate a memory location somewhere in the computer's memory for a variable called width. When drawing this situation, we use a box to represent each memory location. We will use an arrow to keep track of which statements in the program have been executed so far. The arrow will point to the next statement to be executed. A "?" in a box is used to indicate that the value of the variable is unknown, or junk.

```
    int main()
    {
        int width;
        int height;

    -->getDimensions(width, height);
        drawBox(width, height);
    }
    _____

    width        height

     ___          ___
    |   |        |   |
    | ? |        | ? |
    |___|        |___|
```

The next thing that happens is the call to getDimensions. As we enter getDimensions, two parameters are declared, and memory is allocated for these two new variables. If these parameters were pass-by-value parameters, they would get initialized to the values of the corresonding arguments. In this case, the values of the corresponding arguments are unknown, and so the parameters would both get initialized to junk. However, these parameters are not pass-by-value, they are pass-by-reference. The way pass-by-reference works is that instead of being initialized to the value stored in the corresponding argument, the parameters are initialized to the memory address of the corresponding argument. If we assume for a moment that width is stored at memory location 5000 and height is stored at memory location 5004, the situation would look like this:

```
    int main()
    {
        int width;
```

```
    int height;

    getDimensions(width, height);
-->drawBox(width, height);
}

void getDimensions(int &width,
                   int &height)
{
-->cout << "Enter width: ";
    cin >> width;
    cout << "Enter height: ";
    cin >> height;
}
```

```
    width          height

   _____        _____
  |      |      |      |
  |  ?   |      |  ?   |
  |_____|      |_____|



   _____        _____
  | 5000 |      | 5004 |
  |_____|      |_____|

    width          height
```

Computer programmers, however, seldom think in terms of the actual value of memory addresses. A memory address is essentially just a way to refer back to that memory location. Instead of using the actual memory address, we will use an arrow to point from the variable where the memory address is being stored to the memory location itself.

```
int main()
{
    int width;
    int height;

    getDimensions(width, height);
-->drawBox(width, height);
}

void getDimensions(int &width,
                   int &height)
{
-->cout << "Enter width: ";
    cin >> width;
    cout << "Enter height: ";
    cin >> height;
}
```

```
    width          height

   _____        _____
  |      |      |      |
  |  ?   |      |  ?   |
  |__||__|      |__||__|
     /|\           /|\
      |             |
   _____        _____
  |  |   |      |  |   |
  |  |   |      |  |   |
  |_____|      |_____|

    width          height
```

Now we have done the hard part: setting up the parameters. We are ready to start executing the statements inside the function getDimensions. The first statement doesn't interest us here. The second statement does.

Let's say that the user types the value 7 for the width. Where does the 7 go? C++ would normally put it in the local variable width, but when we look in width we see that it is not a normal variable but rather a "reference" to another variable. So instead of putting the 7 in the local variable width, we follow the arrow back to the variable width which belongs to main and put the value 7 there. It is important to realize that the fact that the names are the same is purely coincidental. C++ does not look at the names, rather it simply follows the arrow and uses the variable that the arrow is pointing to. After following the same procedure for the cin >> height; statement, and assuming that the user enters a 4 for the height, our situation is this:

```
    int main()
    {
        int width;
        int height;

        getDimensions(width, height);
  -->drawBox(width, height);
    }

    void getDimensions(int &width,
                       int &height)
    {
        cout << "Enter width: ";
        cin >> width;
        cout << "Enter height: ";
        cin >> height;
  -->
    }
```

```
     width        height

     ___          ___
    |   |        |   |
    | 7 |        | 4 |
    |___|        |___|
    | / | \      | / | \
    |   |        |   |
    |   |        |   |
    |   |        |   |
    |___|        |___|

     width        height
```

Now we have reached the bottom of the getDimensions function. When we reach the bottom of a function, all of the variables declared in that function are deallocated -- for now just think of them as going away. So when we get back to main after having executed the getDimensions function, our situation looks like this:

```
    int main()
    {
        int width;
        int height;

        getDimensions(width, height);
  -->drawBox(width, height);
    }
```

```
     width        height

     ___          ___
    |   |        |   |
    | 7 |        | 4 |
    |___|        |___|
```

The rest of the program involves only pass-by-value. There are no more examples of pass-by-reference. We won't trace through all of the rest of the program, but we will do a bit more for the sake of illustration.

The next statement to be executed is the call to drawBox. Two values are sent. The first value is the value stored in width (7), and the second value is the value stored in height (4). As we enter drawBox, two parameters are declared: width and height. Since width is listed first, it gets initialized to the first value that was sent (7). Since height is listed second, it gets initialized to the second value that was sent (4). Notice that this has nothing to do with the names of the arguments or parameters, but with the order in which they are listed. Here is the picture at this point:

```
    int main()
    {
        int width;
        int height;

        getDimensions(width, height);
        drawBox(width, height);
    -->
    }

    void drawBox(int width, int height)
    {
        drawHorizontalLine(width);
        draw2VerticalLines(width - 2,
                           height - 2);
        drawHorizontalLine(width);
    }
```

```
  width       height

  ┌───┐       ┌───┐
  │ 7 │       │ 4 │
  └───┘       └───┘


  ┌───┐       ┌───┐
  │ 7 │       │ 4 │
  └───┘       └───┘
  width       height
```

Now we are ready to execute the statements inside drawBox. When we call drawHorizontalLine, the value that will be sent is 7, since that is the value stored in width. When we call draw2VerticalLines, the values that will be sent are 5 and 2, the values of width - 2 and height - 2, respectively. And so on.

I have not yet said anything about global variables. Global variables are variables that are declared outside of any function and that can, therefore, be accessed from any function. You should not use global variables, but you should know how they behave in case you have to deal with someone else's code that uses global variables. The rule is this. When a variable name occurs in any function (including main), C++ first looks to see if there is a variable with that name that is local (i.e. declared inside that function). If so, that is the variable that is used. If not, C++ proceeds to look for a global variable with that name. If there is a global variable with that name, it is used. If there is neither a local variable or a global variable with that name, a compile-time error results. One result of this procedure is that if there is a global variable AND a local variable with that name, it is the local version that is used, not the global version

I have provided 6 trace programs. You can find links to these programs below. I strongly suggest that you print these out, try to predict the output, and then compile and run them to see if your answers are correct. You will have to deal with global variables on a few of the trace programs I have provided.

Note about the trace problems: Some of the output from these trace programs will be "junk", since the variables will be uninitialized. Be careful: sometimes this junk will be obviously junk (say you get a result like -80343943). Other times, the junk might be a number like "0", which looks like a real value that was computed, but is really just the value that happened to be in that variable.

trace example 1

Section 7.3: Generating Random Numbers

You will probably need to know how to generate random numbers to do at least one of your assignments in this class. This is done in C++ with the function `rand()`. This function is a value returning function, which means that it is used as an expression when you call it. It returns a random integer between 0 and something very big (roughly 32,000). The following program illustrates the use of the `rand()` function and shows the results when I ran the program. Notice that in order to use the `rand()` function I have to put a `#include <cstdlib>` at the top of my program.

```cpp
#include <iostream>
#include <iomanip>
#include <cstdlib>
using namespace std;

int main()
{
    for (int i = 0; i < 100; i++){
        cout << setw(8) << rand();
        if (i % 5 == 4){
            cout << endl;
        }
    }
}
```

```
16838     5758    10113    17515    31051
 5627    23010     7419    16212     4086
 2749    12767     9084    12060    32225
17543    25089    21183    25137    25566
26966     4978    20495    10311    11367
30054    17031    13145    19882    25736
30524    28505    28394    22102    24851
19067    12754    11653     6561    27096
13628    15188    32085     4143     6967
31406    24165    13403    25562    24834
31353      920    10444    24803     7962
19318     1422    31327    10457     1945
14479    29983    18751     3894    18670
 8259    16248     7757    15629    13306
28606    13990    11738    12516     1414
 5262    17116    22825     3181    13134
25343     8022    11233     7536     9760
 9979    29071     1201    21336    13061
22160    24005    30729     7644    27475
31693    25514    14139    22088    26521
```

Typically we don't want numbers between 0 and 32 thousand something. Let's suppose we want numbers between 0 and 100. So how can we scale these numbers so that we get a result that is between 0 and 100? Let's use the modulus (%) operator. If I change the line

```cpp
cout << setw(8) << rand();
```

to

```cpp
cout << setw(8) << rand() % 101;
```

what will I get? Well, if I take any number % 101, what is the smallest result I could get? 0. What is the largest number I could get? 100. Exactly what I wanted. The following program illustrates this:

```cpp
#include <iostream>
#include <iomanip>
```

```cpp
#include <cstdlib>
using namespace std;

int main()
{
    for (int i = 0; i < 100; i++){
        cout << setw(8) << rand() % 101;
        if (i % 5 == 4){
            cout << endl;
        }
    }
}
```

```
      72         1        13        42        44
      72        83        46        52        46
      22        41        95        41         6
      70        41        74        89        13
     100        29        93         9        55
      57        63        15        86        82
      22        23        13        84         5
      79        28        38        97        28
      94        38        68         2        99
      96        26        71         9        89
      43        11        41        58        84
      27         8        17        54        26
      36        87        66        56        86
      78        88        81        75        75
      23        52        22        93         0
      10        47       100        50         4
      93        43        22        62        64
      81        84        90        25        32
      41        68        25        69         3
      80        62       100        70        59
```

Now we have random numbers between 0 and 100.

There are several other points to keep in mind when working with random numbers.

1. Usually it won't be enough to just print out a random number. You'll want to save the random number in a variable and then use that variable instead of directly printing out the random number. This way the variable will still be available to be used later. So instead of just printing the random number, you might have a statement like `num1 = rand() % 101;`.

2. Notice that the number I put after the % is always ONE BIGGER than the high end of the range of numbers I want to generate. So if I want to generate numbers between 0 and x, I would use a statement like `num1 = rand() % (x + 1);`.

3. What I have told you so far will work fine for generating random numbers, but each time you run your program you will get the same sequence of random numbers. Here's a little trick to make it so you get a different sequence of random numbers each time you run your program. First, put the statement `#include <ctime>` at the top of your program. Then put the statement `srand(static_cast<unsigned>(time(0)));` as the first line in your main function. Make sure the argument is a zero, not an upper case letter 'O'. This function call uses the clock time to initialize the random number generator. Simply put it as the first line in your main function and forget about it. Don't put it anywhere else in your program.